

Time series

Professor McNamara

Time series

What is a time series?

Why model time series?

Example– clothing expenditures


```
## # A tibble: 85 x 3
##   year sales.b     t
##   <dbl>   <dbl> <int>
## 1  1929     9     1
## 2  1930    7.8    2
## 3  1931    6.7    3
## 4  1932    4.9    4
## 5  1933    4.5    5
## 6  1934    5.5    6
## 7  1935    5.8    7
## 8  1936    6.3    8
## 9  1937    6.6    9
## 10 1938    6.5   10
## # ... with 75 more rows
```

Plots

The greatest value of a picture is when it forces us to notice what we never expected to see

-John Tukey

Time series plot



What does this plot tell us?

What model might we fit?

Section 12.1

Variations on linear models

Initially, let's pretend that the only tool we have available to us is linear regression. How can we adapt linear regression to help us model time series data?

Example– clothing expenditures


Using year as predictor,

```
##
## Call:
## lm(formula = sales.b ~ year, data = clothes)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -63.372 -42.288   1.892  35.077  79.259
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -8458.2427   371.6368  -22.76  <2e-16 ***
## year          4.3484     0.1885   23.06  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 42.65 on 83 degrees of freedom
## Multiple R-squared:  0.865, Adjusted R-squared:  0.8634
## F-statistic: 531.9 on 1 and 83 DF, p-value: < 2.2e-16
```

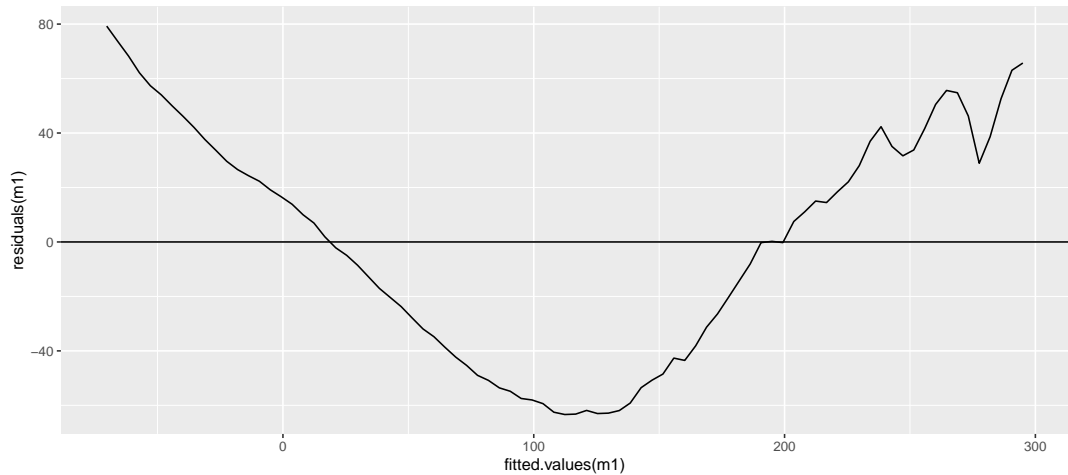
Could also use the time step as the predictor,

```
##
## Call:
## lm(formula = sales.b ~ t, data = clothes)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -63.372 -42.288   1.892  35.077  79.259
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -74.6076     9.3341  -7.993 6.72e-12 ***
## t            4.3484     0.1885  23.064 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

Plotting the fit




Residuals




Quadratic fit?

```
##
## Call:
## lm(formula = sales.b ~ poly(t, degree = 2), data = clothes)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -33.776  -9.731   1.379   9.826  24.538
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      112.372      1.443   77.85  <2e-16 ***
## poly(t, degree = 2)1  983.633     13.307   73.92  <2e-16 ***
## poly(t, degree = 2)2  369.390     13.307   27.76  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.31 on 82 degrees of freedom
## Multiple R-squared:  0.987, Adjusted R-squared:  0.9867
## F-statistic: 3117 on 2 and 82 DF, p-value: < 2.2e-16
```

Quadratic fit?



Residuals



Cosine model

$$Y = \beta_0 + \alpha \cos\left(\frac{2\pi t}{12} + \theta\right) + \epsilon$$

This model is non-linear, but it can be transformed

$$Y = \beta_0 + \alpha \cos(\theta) \cos\left(\frac{2\pi t}{12}\right) - \alpha \sin(\theta) \sin\left(\frac{2\pi t}{12}\right) + \epsilon$$

Then, instead of using t as our predictor, we use $X_{\cos} = \cos\left(\frac{2\pi t}{12}\right)$ and $X_{\sin} = \sin\left(\frac{2\pi t}{12}\right)$

$$Y = \beta_0 + \beta_1 X_{\cos} + \beta_2 X_{\sin} + \epsilon$$


Shiny app

Play with this shiny app to get a sense for what the different coefficients do.

Example– butter imports

```
## # A tibble: 83 x 2
##   month imports_dan
##   <dbl>         <dbl>
## 1     1           NA
## 2     2          166
## 3     3          181
## 4     4          175
## 5     5          187
## 6     6          260
## 7     7          213
## 8     8          199
## 9     9          197
## 10    10          212
## # ... with 73 more rows
```

Example— butter imports




Example– butter imports

```
butter <- butter %>%
  mutate(Xcos = cos(2 * pi * month / 12), Xsin = sin(2 * pi * month / 12))
m4 <- lm(imports_dan ~ Xcos + Xsin, data = butter)
summary(m4)

##
## Call:
## lm(formula = imports_dan ~ Xcos + Xsin, data = butter)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -46.407 -17.930   0.605  16.267  54.219
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   199.05      2.62   75.964 < 2e-16 ***
## Xcos          -33.73      3.74  -9.020 8.91e-14 ***
## Xsin          -12.36      3.67  -3.367 0.00117 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 23.71 on 79 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.5381, Adjusted R-squared:  0.5264
## F-statistic: 46.02 on 2 and 79 DF, p-value: 5.605e-14
```

Example— butter imports



Seasonal means model

A simpler method would be to just use means for each time period.

$$Y = \beta_0 + \beta_1 \textit{Period}_2 + \beta_2 \textit{Period}_3 + \cdots + \beta_{k-1} \textit{Period}_k + \epsilon$$


E.g.,

$$Y = \beta_0 + \beta_1 \textit{Month}_{\textit{February}} + \beta_2 \textit{Month}_{\textit{March}} + \cdots + \beta_{11} \textit{Month}_{\textit{December}} + \epsilon$$

Example– butter imports

```
##
## Call:
## lm(formula = imports_dan ~ as_factor(month_real), data = butter)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -56.857 -15.208  -0.143  14.750  44.714
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      169.833      9.197  18.465 < 2e-16 ***
## as_factor(month_real)2    -3.833     12.534  -0.306 0.760638
## as_factor(month_real)3     8.167     12.534   0.652 0.516818
## as_factor(month_real)4    25.738     12.534   2.053 0.043766 *
## as_factor(month_real)5    54.167     12.534   4.322 5.02e-05 ***
## as_factor(month_real)6    85.024     12.534   6.783 3.09e-09 ***
## as_factor(month_real)7    66.452     12.534   5.302 1.27e-06 ***
## as_factor(month_real)8    45.452     12.534   3.626 0.000542 ***
## as_factor(month_real)9    31.167     12.534   2.487 0.015286 *
## as_factor(month_real)10   27.167     12.534   2.167 0.033605 *
## as_factor(month_real)11    1.310     12.534   0.104 0.917089
## as_factor(month_real)12   13.167     13.007   1.012 0.314899
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 22.53 on 70 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.6306, Adjusted R-squared:  0.5726
## F-statistic: 10.86 on 11 and 70 DF, p-value: 2.338e-11
```

Example— butter imports



Comparing model choices

Cosine trend:

fewer parameters

```
## (Intercept)      Xcos      Xsin
## 199.04722    -33.73367   -12.35982
```

worse R^2

```
## [1] 0.5381284
```

Seasonal means:


better R^2 and adjusted R^2 ,

```
## [1] 0.6306284
## [1] 0.5725843
```

lots of parameters. . .

```
## (Intercept) as_factor(month_real)2 as_factor(month_real)3
## 169.833333 -3.833333 8.166667
## as_factor(month_real)4 as_factor(month_real)5 as_factor(month_real)6
## 25.738095 54.166667 85.023810
## as_factor(month_real)7 as_factor(month_real)8 as_factor(month_real)9
## 66.452381 45.452381 31.166667
## as_factor(month_real)10 as_factor(month_real)11 as_factor(month_real)12
## 27.166667 1.309524 13.166667
```


Another example— butter price




Example– butter price

```
##
## Call:
## lm(formula = price_dan ~ as_factor(month_real), data = butter)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -163.16  -58.51  -23.06   61.86  187.73
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      533.071      32.186   16.562 <2e-16 ***
## as_factor(month_real)2      24.786      45.517    0.545  0.5878
## as_factor(month_real)3     -16.471      45.517   -0.362  0.7185
## as_factor(month_real)4     -67.400      45.517  -1.481  0.1431
## as_factor(month_real)5     -84.200      45.517  -1.850  0.0685 .
## as_factor(month_real)6     -73.186      45.517  -1.608  0.1123
## as_factor(month_real)7     -48.371      45.517  -1.063  0.2915
## as_factor(month_real)8     -24.057      45.517  -0.529  0.5988
## as_factor(month_real)9       2.686      45.517   0.059  0.9531
## as_factor(month_real)10      3.743      45.517   0.082  0.9347
## as_factor(month_real)11      4.614      45.517   0.101  0.9195
## as_factor(month_real)12     12.595      47.376   0.266  0.7911
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 85.15 on 71 degrees of freedom
## Multiple R-squared:  0.169, Adjusted R-squared:  0.04024
## F-statistic: 1.313 on 11 and 71 DF,  p-value: 0.2357
```


Example— butter price



Residuals




Another term


Maybe it would help to include a linear trend,

```
##
## Call:
## lm(formula = price_dan ~ as_factor(month_real) + month, data = butter)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -131.525  -30.432   -1.378   32.357   97.295
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    630.6023     21.5633   29.244 < 2e-16 ***
## as_factor(month_real)2     27.4217     27.7876    0.987  0.32712
## as_factor(month_real)3    -11.1995     27.7908   -0.403  0.68818
## as_factor(month_real)4    -59.4921     27.7959   -2.140  0.03582 *
## as_factor(month_real)5    -73.6561     27.8032   -2.649  0.00997 **
## as_factor(month_real)6    -60.0059     27.8125   -2.158  0.03440 *
## as_factor(month_real)7    -32.5556     27.8239   -1.170  0.24595
## as_factor(month_real)8     -5.6054     27.8374   -0.201  0.84100
## as_factor(month_real)9     23.7735     27.8529    0.854  0.39627
## as_factor(month_real)10    27.4666     27.8705    0.986  0.32777
## as_factor(month_real)11    30.9740     27.8902    1.111  0.27055
## as_factor(month_real)12    25.7751     28.9461    0.890  0.37627
## month           -2.6360         0.2401  -10.978 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 51.98 on 70 degrees of freedom
## Multiple R-squared:  0.6947, Adjusted R-squared:  0.6423
## F-statistic: 13.27 on 12 and 70 DF,  p-value: 1.05e-13
```

Another term







Residuals



Decomposing time series

It is useful to think of time series as a decomposition



Section 12.2





Lags and autocorrelation

Random walk


In a random walk, each value is just a random movement from the previous one.

$$Y_t = Y_{t-1} + \epsilon_t$$

Random walks



Example— Apple Stock



Differences

First difference of a time series,

$$\Delta Y_t = Y_t - Y_{t-1}$$

For a random walk with $Y_t = Y_{t-1} + \epsilon_t$,

$$\Delta Y_t = Y_t - Y_{t-1} = \epsilon_t$$


Example— Apple Stock

```
AppleStock <- AppleStock %>%  
  tibble() %>%  
  mutate(diff = Price - lag(Price))  
AppleStock %>%  
  select(Price, diff, day)
```

```
## # A tibble: 66 x 3  
##   Price  diff day  
##   <dbl> <dbl> <date>  
## 1  99.4 NA    2016-07-21  
## 2  98.7 -0.77 2016-07-22  
## 3  97.3 -1.32 2016-07-25  
## 4  96.7 -0.67 2016-07-26  
## 5 103.   6.28 2016-07-27  
## 6 104.   1.39 2016-07-28  
## 7 104.  -0.13 2016-07-29  
## 8 106.   1.84 2016-08-01  
## 9 104.  -1.57 2016-08-02  
## 10 106.   1.31 2016-08-03  
## # ... with 56 more rows
```

First differences

Warning: Removed 1 row(s) containing missing values (geom_path).



Autocorrelation

You can measure the association between time series values that are k time units apart. For example, could also do lag 2.

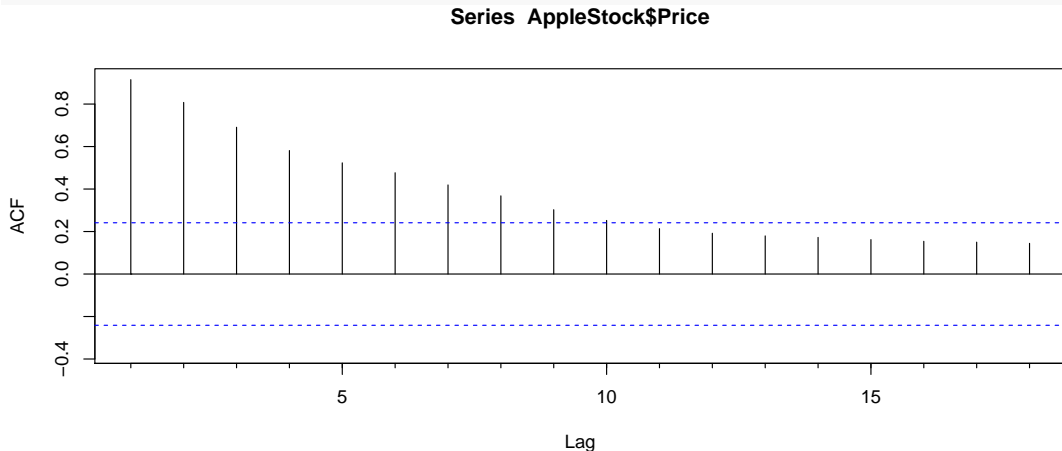
```
AppleStock <- AppleStock %>%  
  mutate(diff2 = Price - lag(Price, n = 2))  
AppleStock %>%  
  select(Price, diff2, day)
```

```
## # A tibble: 66 x 3  
##   Price  diff2 day  
##   <dbl> <dbl> <date>  
## 1  99.4  NA    2016-07-21  
## 2  98.7  NA    2016-07-22  
## 3  97.3 -2.09  2016-07-25  
## 4  96.7 -1.99  2016-07-26  
## 5 103.   5.61  2016-07-27  
## 6 104.   7.67  2016-07-28  
## 7 104.   1.26  2016-07-29  
## 8 106.   1.71  2016-08-01  
## 9 104.   0.27  2016-08-02  
## 10 106.  -0.260  2016-08-03  
## # ... with 56 more rows
```

Computing autocorrelation

One way to compute autocorrelation is with `forecast::Acf()`

```
Acf(AppleStock$Price)
```




```
Acf(AppleStock$Price, plot = FALSE)
```

```
##  
## Autocorrelations of series 'AppleStock$Price', by lag  
##  
##    0    1    2    3    4    5    6    7    8    9   10   11   12  
## 1.000 0.915 0.808 0.691 0.581 0.523 0.477 0.419 0.368 0.303 0.252 0.214 0.192  
##   13   14   15   16   17   18
```

No autocorrelation

Let's compare those last few plots with an autocorrelation plot of one of the lags

Series AppleStock\$diff



Nothing makes it to the edge, so we're good!

Stationarity


One reason to look at the ACF plot is to determine if a time series is “stationary.” Stationarity basically means things are staying the same over time.

Ways to see **lack of stationarity**:

- overall increasing or decreasing trend in data plot
- linear trend in ACF plot


Example– Apple prices

Apple stock prices are **not** stationary. We can see this either in the plot of the data,



Or in the ACF plot,



Series AppleStock\$Price



Example— Apple stock price differences

The first differences of Apple stock prices **are** stationary

Warning: Removed 1 row(s) containing missing values (geom_path).



Seasonal differences


We can also have differences on regular lags, such as seasons. For example, we could find

$$\Delta_{12} Y_t = Y_t - Y_{t-12}$$

to compare to the same month from the previous year


Back to butter prices

Recall our butter prices data,





We could look at an ACF plot for this data,

Series butter\$price_dan



Butter price differences

Warning: Removed 1 row(s) containing missing values (geom_path).





Section 12.3

ARIMA models

AutoRegressive (AR)

Moving Average (MA)



Example— clothing expenditures



Percent difference

Instead of just subtracting the previous month to get an absolute difference, let's do a percent difference.

```
## Warning: Removed 1 row(s) containing missing values (geom_path).
```



Autoregressive model

ϕ

$$Y_t = \delta + \phi Y_{t-1} + \epsilon_t$$

Example– clothing expenditures

Find the AR(1) model

```
Arima(clothes$pctdiff, order = c(1, 0, 0), include.constant = TRUE)

## Series: clothes$pctdiff
## ARIMA(1,0,0) with non-zero mean
##
## Coefficients:
##          ar1      mean
##          0.6196  0.0366
## s.e.      0.0903  0.0160
##
## sigma^2 estimated as 0.003281:  log likelihood=121.81
## AIC=-237.62   AICc=-237.32   BIC=-230.32
```

Annoyingly in R, the “mean” isn’t the constant term in the model. You have to do a little algebra,

$$\mu = \delta + \phi_1 \mu$$

Considering more history

We don't have to just use the previous value! We could fit $AR(p)$, with as many lags as we want.

$$Y_t = \delta + \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \cdots + \phi_p Y_{t-p} + \epsilon_t$$

Example– clothing expenditures

Fitting the AR(2) model,

```
Arima(clothes$pctdiff, order = c(2, 0, 0), include.constant = TRUE)
```

```
## Series: clothes$pctdiff
## ARIMA(2,0,0) with non-zero mean
##
## Coefficients:
##          ar1          ar2          mean
##          0.6614   -0.0787   0.0376
## s.e.    0.1090    0.1150   0.0147
##
## sigma^2 estimated as 0.003303:  log likelihood=122.04
## AIC=-236.08   AICc=-235.58   BIC=-226.36
```

How do we know which variables are significant?

```
Arima(clothes$pctdiff, order = c(2, 0, 0), include.constant = TRUE)
```

```
## Series: clothes$pctdiff
## ARIMA(2,0,0) with non-zero mean
##
## Coefficients:
##          ar1      ar2      mean
##          0.6614 -0.0787  0.0376
## s.e.    0.1090   0.1150  0.0147
##
## sigma^2 estimated as 0.003303:  log likelihood=122.04
## AIC=-236.08   AICc=-235.58   BIC=-226.36
```

Rule of thumb:

$$|\hat{\phi}_i / SE| > 2$$

More formal

Can use `lmtest::coefstest()`

```
##
## z test of coefficients:
##
##           Estimate Std. Error z value Pr(>|z|)
## ar1          0.661368   0.108984   6.0685 1.291e-09 ***
## ar2         -0.078663   0.115013  -0.6839  0.49401
## intercept    0.037602   0.014654   2.5660  0.01029 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Moving Average (MA)

Okay, so we've talked about AR models, now let's talk MA models. These models consider that values in a time series might be related to the residual(s) from previous time steps.

$$Y_t = \delta + \epsilon_t + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \cdots + \theta_q \epsilon_{t-q}$$

Note: often this type of model would have minus signs holding it together, rather than plus signs. But, the `Arima()` function does things a little different, so we're being consistent with that.

Example– clothing expenditures

```
(MA1 <- Arima(clothes$spctdiff, order = c(0, 0, 1), include.constant = TRUE))
```

```
## Series: clothes$spctdiff
## ARIMA(0,0,1) with non-zero mean
##
## Coefficients:
##          ma1      mean
##          0.7429  0.0388
## s.e.    0.0963  0.0107
##
## sigma^2 estimated as 0.003294:  log likelihood=121.47
## AIC=-236.94   AICc=-236.64   BIC=-229.65
```

```
coeftest(MA1)
```

```
##
## z test of coefficients:
##
##          Estimate Std. Error z value  Pr(>|z|)
## ma1          0.742912   0.096256   7.7181 1.181e-14 ***
## intercept 0.038788    0.010735   3.6131 0.0003026 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

ARIMA models

We can put the AR and MA models together,

$$\begin{aligned} Y_t = & \delta + \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \cdots + \phi_p Y_{t-p} \epsilon_t + \epsilon_t \\ & + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \cdots + \theta_q \epsilon_{t-q} \end{aligned}$$

Example– clothing expenditures

```
(ARMA11 <- Arima(clothes$pctdiff, order = c(1, 0, 1), include.constant = TRUE))
```

```
## Series: clothes$pctdiff
## ARIMA(1,0,1) with non-zero mean
##
## Coefficients:
##          ar1      ma1      mean
##      0.4817  0.2077  0.0377
## s.e.  0.2185  0.2744  0.0142
##
## sigma^2 estimated as 0.003292:  log likelihood=122.16
## AIC=-236.33  AICc=-235.82  BIC=-226.6
```

```
coeftest(ARMA11)
```

```
##
## z test of coefficients:
##
##      Estimate Std. Error z value Pr(>|z|)
## ar1      0.481688   0.218519   2.2043 0.027501 *
## ma1      0.207720   0.274355   0.7571 0.448977
## intercept 0.037708   0.014207   2.6542 0.007949 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

One more parameter

We can also bring differences into the model,

$$\begin{aligned}\Delta^d Y_t = \delta &+ \phi_1 \Delta^d Y_{t-1} + \cdots + \phi_p \Delta^d Y_{t-p} \epsilon_t + \epsilon_t \\ &+ \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \cdots + \theta_q \epsilon_{t-q}\end{aligned}$$

Generally, we either do $d = 0$ or $d = 1$

ARIMA for clothing expenditure

Since we're going to include differences with the d term, let's go back to original sales

```
(ARIMA11 <- Arima(clothes$sales.b, order = c(1, 1, 0), include.constant = TRUE))
```


```
## Series: clothes$sales.b
## ARIMA(1,1,0) with drift
##
## Coefficients:
##          ar1    drift
##          0.6191  4.1390
## s.e.      0.0851  1.1233
##
## sigma^2 estimated as 16.36:  log likelihood=-235.81
## AIC=477.62   AICc=477.92   BIC=484.91
```

```
coeftest(ARIMA11)
```

```
##
## z test of coefficients:
##
##      Estimate Std. Error z value Pr(>|z|)
## ar1    0.619071   0.085051  7.2788 3.368e-13 ***
## drift  4.138951   1.123329  3.6845 0.0002291 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual ACF

We want to make sure the residuals look random, to ensure we've modeled out any signal in the noise.



Still some seasonality there

Seasonal ARIMA

Let's just go nuts with terms. . . This model includes:

- p regular autoregressive terms

- d regular differences

- q moving average terms

- P seasonal autoregressive terms

- D seasonal differences

- Q seasonal moving average terms

With this much going on, we're just going to let R handle it

Seasonal ARIMA


```
(ARIMA112 <- Arima(clothes$sales.b, order = c(1, 1, 0), seasonal = list(order= c(2,0,0), period=12)))
```

```
## Series: clothes$sales.b
## ARIMA(1,1,0)(2,0,0)[12]
##
## Coefficients:
##          ar1      sar1      sar2
##          0.6499  0.3552  0.1013
## s.e.      0.0995  0.1574  0.1531
##
## sigma^2 estimated as 16.65:  log likelihood=-237.21
## AIC=482.41   AICc=482.92   BIC=492.14
```

```
coeftest(ARIMA112)
```

```
##
## z test of coefficients:
##
##      Estimate Std. Error z value Pr(>|z|)
## ar1   0.64985    0.09948  6.5325 6.468e-11 ***
## sar1   0.35519    0.15736  2.2571  0.0240 *
## sar2   0.10131    0.15315  0.6615  0.5083
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


Residuals



Forecasting

It's tough to make predictions, especially about the future
-Yogi Berra

Forecasting

```
ARIMA111 <- Arima(clothes$sales.b, order = c(1, 1, 0), seasonal = list(order= c(1,0,0), period=12))  
forecast(ARIMA111, h=10)
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## 86	366.4090	361.1859	371.6320	358.4210	374.3969
## 87	372.4280	362.2929	382.5632	356.9277	387.9284
## 88	379.4049	364.4613	394.3484	356.5507	402.2591
## 89	385.9650	366.4806	405.4493	356.1663	415.7636
## 90	390.5975	366.8864	414.3085	354.3345	426.8604
## 91	392.6096	364.9815	420.2378	350.3560	434.8633
## 92	391.4758	360.2158	422.7358	343.6677	439.2839
## 93	386.8578	352.2204	421.4952	333.8845	439.8311
## 94	392.3946	354.6036	430.1856	334.5983	450.1910
## 95	399.4535	358.7043	440.2026	337.1330	461.7739

Plotting forecasts

```
pred1 <- forecast(ARIMA111, h=10)  
plot(pred1)
```

Forecasts from ARIMA(1,1,0)(1,0,0)[12]

